



Department of Mathematics and Computer Science
Database Group

Efficient Graph Processing Using AWS S3

Master Thesis

Aditya Chandla

Supervisors:
dr. Nikolay Yakovets

Eindhoven, July 2024

Abstract

As the size of graphs being queried by graph databases increases, coupling storage and compute increases the cost and limits the flexibility of traditional graph database management systems (GDBMS). As part of this thesis, we explore the viability of an architecture where the compute and storage are independent. This independence is achieved using a distributed cloud storage service (AWS S3) which provides bottomless storage and theoretically unlimited throughput. Using this distributed cloud storage solution, we evaluate the latency of running two common graph traversal algorithms i.e breadth first search (BFS) and depth first search (DFS). We then compare this latency with some other systems which may be used to perform BFS and DFS on large graphs.

Preface

Please write all your preface text here. If you do so, don't forget to thank your supervisor, other committee members, your family, colleagues etc. etc.

Contents

Contents	vii
List of Figures	ix
List of Tables	xi
Listings	xiii
1 Introduction	1
1.1 Structure of the thesis	1
2 Preliminaries	3
2.1 Distributed storage services	3
2.2 Serverless Architectures	3
2.3 Caching and Prefetching in distributed systems	3
3 System Architecture	5
3.1 Component Overview	5
3.2 Query Evaluation	5
3.3 Baseline Implementation	5
3.4 Graph access service	5
3.4.1 Modified CSR structure	5
3.4.2 Caching	5
3.4.3 Prefetching	5
3.5 Parallelizing graph algorithm service	5
3.5.1 Parallel BFS implementation	5
3.5.2 Parallel DFS implementation	5
3.6 Optimizing partition sizes	5
4 Evaluation	7
4.1 Comparison with baseline	7
4.2 Comparison with other tools	7
4.2.1 Neo4j	7
4.2.2 Apache Flink	7
5 Discussion	9
5.1 Alternate system architectures	9
5.2 Use cases for the system	9
5.3 Threats to credibility of this work	9
6 Conclusion and Future Work	11
6.1 Future Work	11

Bibliography	13
Appendix	15
A My First Appendix	15

List of Figures

List of Tables

Listings

Chapter 1

Introduction

In a recent survey titled ‘Ubiquity of large graphs’[6], the authors noted that the size of graphs being used in both academia and industry is increasing. Although this survey found that it was common for graphs to have tens of billions of edges, it also found that scalability was the primary concern of all the participants surveyed[6]. In this thesis, we evaluate one way to alleviate this concern of scalability by decoupling storage of base graph from the query evaluation.

Today, all major cloud providers have a distributed cloud storage offering, Google cloud platform offers Google cloud storage[5], Azure offers Azure Blob storage[2], and AWS offers Simple Storage Service or S3[1]. These distributed storage services are already used by Big Data Analytics tools like Snowflake[7] and RDBMSs like Neon[4] for storage of underlying data. The use of these distributed storage services has not yet been explored for storing and querying graph data.

In this paper, we provide an initial analysis of how distributed storage services can be used to query large graphs. While using these distributed storage services, the primary issue that needs to be addressed is of latency. The latency of accessing data in a networked distributed system involves communication over the network which is at least three orders of magnitude more than accessing data from local storage. In return for this increased latency, we get almost unlimited throughput as read operations are distributed across a cluster of thousands of physical machines if not more. Therefore, in this paper, we evaluate ways to reduce the latency of accessing graphs. More precisely, we will focus on the performance of two commonly used graph traversal algorithms: Breadth first search (BFS) and Depth first search (DFS). We use these algorithms on LDBC datasets to provide results for multi-hop traversals.

1.1 Structure of the thesis

In Chapter-2, we provide the necessary background for this thesis. In Section-2.1 we talk about the history, architecture, and characteristics of distributed storage systems. We also give a brief overview of the capabilities of AWS S3, which is the service that we use in this thesis. Then, in Section-2.2, we discuss the serverless architecture of databases and its advantages over traditional architectures. Finally, in Section-2.3 we discuss the caching and prefetching strategies that we employ in this thesis to lower the latency of graph traversals.

In Chapter-3, we introduce our architecture for performing traversals. We start by providing a description of each component and their responsibilities in Section-3.1. After that, we describe the traversal queries which will be evaluated by our system in Section-3.2. Then we provide a baseline implementation which will be useful for comparing the impact of the improvements made in the subsequent sections. Finally, in Sections-3.4, 3.5, and 3.6, we describe the details of our proposed architecture which enables low latency traversals for large graphs.

In Chapter-4, we present the results of the implementation of the proposed system architecture. We begin this chapter by comparing the performance with the baseline solution in Section-4.1.

Then in Section-4.2, we compare the performance of our solution with Neo4j and Apache Flink. This section highlights various characteristics of different types of tools(GDBMS, RDBMS, Big Data tools, and Custom Solutions) and the areas in which they are suitable.

In Chapter-5, we elucidate the reasoning for various choices made throughout the thesis. In Section-5.1 we discuss other possible architectures for performing traversals on large graphs using S3 and their advantages and disadvantages over the proposed architecture. Then, in Section-5.2, we discuss the use cases where this architecture can be more cost efficient and flexible compared to other tools and where users would be better off avoiding this architecture. Finally, in Section-5.3, we consider the threats to the credibility of this work.

Finally, in Chapter-6, we conclude the thesis and suggest possible directions for future work. This section contains information about how we may be able to extend this work to reach a point where we have a fully functioning graph database whose storage resides in S3.

Chapter 2

Preliminaries

2.1 Distributed storage services

2.2 Serverless Architectures

2.3 Caching and Prefetching in distributed systems

Chapter 3

System Architecture

3.1 Component Overview

3.2 Query Evaluation

3.3 Baseline Implementation

3.4 Graph access service

3.4.1 Modified CSR structure

3.4.2 Caching

3.4.3 Prefetching

3.5 Parallelizing graph algorithm service

3.5.1 Parallel BFS implementation

3.5.2 Parallel DFS implementation

3.6 Optimizing partition sizes

Chapter 4

Evaluation

4.1 Comparison with baseline

4.2 Comparison with other tools

4.2.1 Neo4j

Monolithic Deployment

Distributed Deployment

4.2.2 Apache Flink

Chapter 5

Discussion

5.1 Alternate system architectures

5.2 Use cases for the system

5.3 Threats to credibility of this work

Chapter 6

Conclusion and Future Work

6.1 Future Work

Bibliography

- [1] AWS. Aws s3. <https://aws.amazon.com/s3/>. 1
- [2] Azure. Azure blob storage. <https://azure.microsoft.com/en-us/products/storage/blobs>. 1
- [3] Apache Software Foundation. Apache flink. <https://flink.apache.org/>.
- [4] Neon. Neon. <https://neon.tech/>. 1
- [5] Google Cloud Platform. Google cloud storage. <https://cloud.google.com/storage>. 1
- [6] Siddhartha Sahu, Amine Mhedhbi, Semih Salihoglu, Jimmy Lin, and M Tamer Özsu. The ubiquity of large graphs and surprising challenges of graph processing. *Proceedings of the VLDB Endowment*, 11(4):420–431, 2017. 1
- [7] Snowflake. Snowflake. <https://www.snowflake.com/en/>. 1

Appendix A

My First Appendix

In this file (`appendices/main.tex`) you can add appendix chapters, just as you did in the `thesis.tex` file for the ‘normal’ chapters. You can also choose to include everything in this single file, whatever you prefer.