



Department of Mathematics and Computer Science
Database Group

Efficient Graph Processing Using AWS S3

Master Thesis

Aditya Chandla

Supervisors:
dr. Nikolay Yakovets

Eindhoven, July 2024

Abstract

As the size of graphs being queried by graph databases increases, coupling storage and compute increases the cost and limits the flexibility of traditional graph database management systems (GDBMS). As part of this thesis, we explore the viability of an architecture where the compute and storage are independent. This independence is achieved using a distributed cloud storage service (AWS S3) which provides bottomless storage and theoretically unlimited throughput. Using this distributed cloud storage solution, we evaluate the latency of running two common graph traversal algorithms i.e breadth first search (BFS) and depth first search (DFS). We then compare this latency with some other systems which may be used to perform BFS and DFS on large graphs.

Preface

Please write all your preface text here. If you do so, don't forget to thank your supervisor, other committee members, your family, colleagues etc. etc.

Contents

Contents	iv
1 Introduction	1
1.1 Research Objectives	1
1.2 Structure of the thesis	2
2 Preliminaries	3
2.1 Distributed storage services	3
2.2 Separating compute and storage	5
2.3 Caching and Prefetching	6
3 System Architecture	9
3.1 Component Overview	9
3.2 Baseline Implementation	10
3.3 Graph access service	10
3.3.1 Modified CSR structure	10
3.3.2 Caching and Prefetching	11
3.4 Parallelizing graph algorithm service	13
3.4.1 Parallel BFS implementation	13
3.4.2 Parallel DFS implementation	13
4 Evaluation	15
4.1 Comparison with baseline	15
4.1.1 Optimizing partition sizes	15
4.2 Comparison with other tools	15
4.2.1 Neo4j	15
4.2.2 Apache Flink	15
5 Discussion	17
5.1 Alternate system architectures	17
5.2 Use cases for the system	17
5.3 Threats to credibility of this work	17
6 Conclusion and Future Work	19
6.1 Future Work	19
Bibliography	21
Appendix	23
A System Details	23
A.1 Query Evaluation	23

Chapter 1

Introduction

In a recent survey titled ‘Ubiquity of large graphs’[16], the authors noted that the size of graphs being used in both academia and industry is increasing. Although this survey found that it was common for graphs to have tens of billions of edges, it also found that scalability was the primary concern of all the participants surveyed[16]. In this thesis, we evaluate one way to alleviate this concern of scalability by decoupling storage of base graph from the query evaluation.

Today, all major cloud providers have a distributed cloud storage offering, Google cloud platform offers Google cloud storage[15], Azure offers Azure Blob storage[4], and AWS offers Simple Storage Service or S3[3]. These distributed storage services are already used by Big Data Analytics tools like Snowflake[18] and RDBMSs like Neon[14] for storage of underlying data. The use of these distributed storage services has not yet been explored for storing and querying graph data.

1.1 Research Objectives

In this paper, we provide an initial analysis of how distributed storage services can be used to query large graphs. While using these distributed storage services, the primary issue that needs to be addressed is of latency. The latency of accessing data in a networked distributed system involves communication over the network which is at least three orders of magnitude more than accessing data from local storage. In return for this increased latency, we get virtually unlimited throughput as read operations are distributed across a cluster of thousands of physical machines if not more. Therefore, in this paper, we evaluate ways to reduce the latency of accessing graphs. We will focus on the performance of two commonly used graph traversal algorithms: Breadth first search (BFS) and Depth first search (DFS). More formally, the first research objective is as follows:

RO 1: Gauge the effectiveness of caching and prefetching techniques on the latency of graph access for graph traversals.

Apart from reducing the latency, we will also discuss how the cost model of distributed storage engines is different from the traditional model of coupled compute and storage. For storage engines like S3, we are charged for every gigabyte of storage and we are charged for every request that we make on that storage. On the other hand, in case of an SSD/HDD storage unit, you pay for a certain amount of storage and there is no additional cost for accessing the storage. Therefore, we seek to provide a model to help users choose between one form of storage over the other. More formally, the second research objective is as follows:

RO 2: Provide a cost model to help users decide whether using S3 instead of SSD/HDDs might be more cost effective.

1.2 Structure of the thesis

In Chapter-2, we provide the necessary background for this thesis. In Section-2.1 we talk about the history, architecture, and characteristics of distributed storage systems. We also give a brief overview of the capabilities of AWS S3, which is the service that we use in this thesis. Then, in Section-2.2, we discuss the advantages of database architectures which separate compute and storage. Finally, in Section-2.3 we discuss the caching and prefetching strategies that we employ in this thesis to lower the latency of graph traversals.

In Chapter-3, we introduce our architecture for performing traversals. We start by providing a description of each component and their responsibilities in Section-3.1. Then we provide a baseline implementation which will be useful for comparing the impact of the improvements made in the subsequent sections. Finally, in Sections-3.3 and 3.4 we describe the details of our proposed architecture which enables low latency traversals for large graphs.

In Chapter-4, we present the results of the implementation of the proposed system architecture. We begin this chapter by comparing the performance with the baseline solution in Section-4.1. Then in Section-4.2, we compare the performance of our solution with Neo4j and Apache Flink. This section highlights various characteristics of different types of tools (GDBMS, RDBMS, Big Data tools, and Custom Solutions) and the areas in which they are suitable.

In Chapter-5, we elucidate the reasoning for various choices made throughout the thesis. In Section-5.1 we discuss other possible architectures for performing traversals on large graphs using S3 and their advantages and disadvantages over the proposed architecture. Then, in Section-5.2, we discuss the use cases where this architecture can be more cost efficient and flexible compared to other tools and where users would be better off avoiding this architecture. Finally, in Section-5.3, we consider the threats to the credibility of this work.

Finally, in Chapter-6, we conclude the thesis and suggest possible directions for future work. This section contains information about how we may be able to extend this work to reach a point where we have a fully functioning graph database whose storage resides in S3.

Chapter 2

Preliminaries

The purpose of this thesis is to reduce the latency of graph access in cases where the graph is stored on a distributed file system like AWS S3. In order to put this work into context, we start by providing an overview of distributed storage services and more details about the particular storage service which will be used in this thesis. Then, we provide background on the idea of separating compute and storage, the evolution of this idea, and its applicability. Finally, we describe the current literature on how others have approached the idea of reducing access latency by using caching algorithms.

2.1 Distributed storage services

The idea of accessing files via the network can be traced back to the early days of the internet. However, the first widely used implementation of a networked filesystem was developed by SUN Microsystems[13]. Their implementation is widely referred to as the networked filesystem(NFS). Their implementation provided users with the ability to mount a filesystem present on a remote machine. However, this system did not provide any distribution transparency as the users had to be aware of where and how each file is stored on remote machines.

The first widely used implementation of a distributed filesystem is considered to be the Hadoop Distributed Filesystem (HDFS)[17]. This implementation provided distribution transparency and consistency guarantees on various operations on the files stored in HDFS. This implementation was inspired by the Google Filesystem[9] which was developed by Google in 2003 almost seven years before HDFS.

Although filesystems like HDFS were widely used, filesystems like AWS S3[3] provided an additional feature of being multi-tenant. In other words, with filesystems like AWS S3, a user could reap the benefits of a distributed filesystem without having to manage or pay for an entire cluster. The cost of managing, scaling, and maintaining the distributed cluster was delegated to cloud providers like AWS. Now, users could pay for the amount of data that they store, and the number of requests that they make to their data. This has enabled the usage of such distributed filesystems in applications like video streaming, database storage[18], web content delivery, and backup storage.

The exact architecture of AWS S3 or any similar service provided by competing cloud providers is not known, however, we can glean some information about S3 based on how other distributed filesystems were designed. We will explain some basic ideas related to AWS S3 using Figure-2.1. Let us assume we need to store four files in AWS S3. In order to do that, we first create a storage unit called a 'bucket' which is similar to a directory in a filesystem. After uploading our files to this bucket, the files get replicated over the entire AWS S3 cluster in that particular region. Figure-2.1 assumes that this cluster has a replication factor of two and therefore every file is stored on two physical servers. Now, if users send requests to access these files, their requests can be redirected

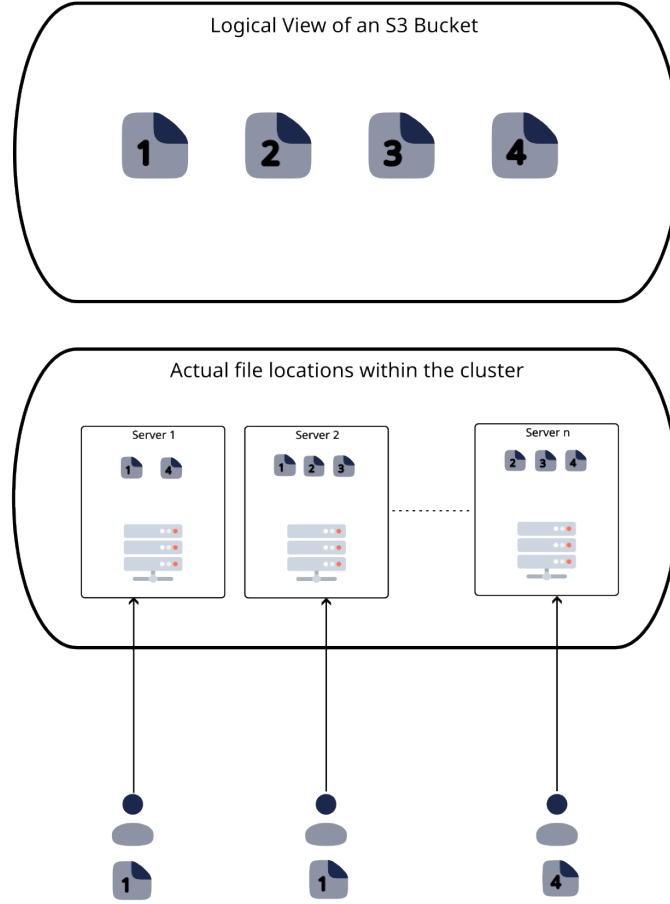


Figure 2.1: Distributed File Store

to any one of the servers containing the desired file. For example, if the first and second user both want to access file number 1, their requests can be served by two different servers since the same copy will be saved on two different servers. For simplicity, the figure shows that users make direct requests to the physical machines containing these files, however, in reality they all send requests to a single endpoint and their requests are then redirected to a server based on some load balancing scheme. In later chapters, we will use this abstract model of a distributed filesystem to argue about the applicability of various techniques.

Apart from the basic understanding of AWS S3 as it related to a distributed filesystem, we also want to draw the reader's attention to the following characteristics of AWS S3 as mentioned in its documentation:

1. **Throughput SLA:** At the time of writing this paper, AWS gaurantees that S3 can handle at least 5500 requests per object stored in a bucket. Note that there is no restriction on the total number of objects that can be stored in a bucket.
2. **Storage Classes:** S3 offers various storage classes that have different cost and performance characteristics. These storage classes determine how much you pay for storage and for making requests. In this thesis, we make use of the 'S3 Standard' and 'Express One Zone' storage classes.
3. **Costs:** The cost of using AWS S3 depends on two factors: total storage and number of

requests. The storage cost for ‘S3 Standard’ storage class is \$0.023/GB and the cost of GET requests is \$0.0004/1k requests. However, the storage cost for the lower latency ‘Express One Zone’ storage type is higher with \$0.16/GB but the cost of making requests is lower at \$0.0002/1k requests.

4. **Artificial subdirectories:** Although AWS CLI and other tools provide the capability to upload and navigate a bucket like a directory, but the underlying architecture only supports having files inside a bucket. Therefore, if you have a folder ‘f1’ with two files ‘one’ and ‘two’, unlike a filesystem, the bucket would only contain two files namely: ‘f1/one’ and ‘f1/two’. The name of the folder is simply prepended to the file names.
5. **Hashing object names:** In order to decide which partitions an object is assigned to, S3 internally hashes the first few characters of the filename. Therefore, if we have object names whose first few characters are the same, it is possible they will all be assigned to a single partition which may lead to lower performance. Although, it is mentioned that S3 automatically splits partitions in case of overload, there is some overhead and time delay in performing the split.
6. **Consistency:** Despite being a distributed system, AWS S3 provides a strong ‘Read-After-Write’ consistency. This means that any updates to an object are immediately visible to all clients.

2.2 Separating compute and storage

One of the first proposed implementation of a database with separate compute and storage can be traced back to 2008[6], where Brantner et al. proposed an architecture of a relational database which used AWS S3 as its storage layer. The authors of this paper primarily focused on how to handle updates under various consistency models when the data is stored in S3. This idea was further developed over the years by AWS and culminated with the release of AWS Aurora[20]. Although Aurora separated compute and storage, the write throughput was still limited due to the fact that all write traffic was directed to a single instance. The first database to support multiple read and write instances on top of distributed data is considered to be Snowflake[7]. Their architecture was built upon the storage layer of AWS S3 on top of which users could create as many ‘virtual warehouses’ as they needed. Since then, various databases with independent storage and compute planes like Neon[14] and AWS Aurora Serverless[2] have been released.

The term ‘Serverless databases’ is often used for databases with separate compute and storage, however, these two concepts are different. A ‘serverless database’ has a pay-as-you-go model where you pay for the amount of storage that your databases uses and the compute resources that were used in a fixed amount of time. In this model, the users simply have to define the minimum and maximum compute resources that the database can use and the databases scales according to the load. Note that the definition of serverless database does not necessitate separation between compute and storage and there are indeed databases like CockroachDB[19] which offer a serverless model although the underlying architecture couples storage and compute together. Therefore, although separating compute and storage may be amenable to a serverless design, it is not a requirement for it. The term serverless, therefore, defines a type of payment model whereas having separate storage and compute for a database is more of a system architecture choice.

A database which separates compute and storage offers better scalability, elasticity, and modularity. When the compute and storage are separated, both of them can be scaled independently. A heavily used service which has relatively low amount of data can scale the database’s compute without wasting any storage. On the other hand, a service which has a lot of data but has a low query workload can scale the database’s storage without wasting any compute. Furthermore, this separation also provides more elasticity so that the database can seamlessly scale up when the workload increases and can scale down during low traffic hours. Finally, separating storage and

compute can lead to more modularity since the storage layer and the compute layer can be upgraded and changed based on the system's requirements. For example, if a cloud provider launches a new instance type which is better suited for your workload, you can switch to the newer instance type without being concerned about data migration. These aforementioned characteristics make databases with separate compute and storage attractive to a variety of use-cases.

Separating compute and storage in databases often comes at the cost of increased latency. This latency is a result of the network communication between the storage layer and compute layer that needs to take place to complete any database operation. Although the throughput of networks has massively increased in the past few years (AWS offers instances with upto 200 Gbps bandwidth), the latency of network communication is limited by the speed of light. Therefore, a database with a local SSD containing all the base data will always be faster than a database whose storage needs to be accessed over a network. The actual difference between performance depends on the physical distance between storage and compute nodes. If the entire deployment is in a single datacenter, the latency is around $500\mu\text{s}$ which is still around 10 times higher than access latency for an SSD[10]. This advantage becomes less pronounced in distributed databases with coupled storage and compute since they may also require some communication between instances in order to respond to user queries. Although there are various caching and prefetching mechanisms to reduce latency for databases with separate storage and compute, there still remain cases where coupled storage and compute yields lower latencies.

2.3 Caching and Prefetching

The concept of caching and prefetching in case of graphs enable us to take advantage of spatial and temporal locality to reduce the latency of future graph accesses. The terms spatial and temporal locality come from literature on processor caches. In the context of a processor cache, temporal locality refers to the tendency of a memory location that is accessed now to be accessed again in the near future. Similarly, spatial locality refers to the tendency of a memory location close to a recently accessed memory location to be accessed in the near future. These concepts can be translated to graphs as follows:

1. **Spatial Locality** in case of graphs means that if a node is accessed, then it is likely that this node's neighbours will be accessed in the near future.
2. **Temporal Locality** in case of graphs refers to the existence of nodes that are central according to some centrality metric. There are various centrality measures[11] which may help us understand which nodes are central for a particular graph algorithm. The main idea is that there exist some nodes which have higher probability of being accessed while running a particular algorithm on a graph.

In this sections, we will discuss some background research on caching and prefetching which would help us exploit spatial and temporal locality to lower the data access latency for graph traversals.

Any caching scheme consists of two fundamental operations: admission of data to the cache and eviction of data from the cache. There are various cache eviction policies like Least Recently Used (LRU), Least Frequently Used (LFU) which define ways to find out what data to evict from the cache when it becomes full. There also exist eviction policies which rely on metadata related to the data in order to decide which data elements to evict. Such metadata might include size of the data element, time of last data access, and access latency of the element in case of a cache miss. Apart from eviction, a cache also needs an admission policy which, in case of most caches, is simply to add a piece of data in case of a cache miss. However, there do exist more sophisticated admission algorithms like TinyLFU[8] which make the admission decision based on metadata related to a data object, which in case of TinyLFU is the access frequency of a data item. Storing metadata related to an object often results in increased memory and maintainance overhead. These admission and eviction policies determine the effectiveness of a cache algorithm for a given use-case.

We will first consider schemes which do not require additional information about the data being stored like size or access latency in case of a cache miss. In this case, LRU and LFU are two of the most commonly used algorithms. It has been shown that if the access pattern follows Zipf's law[21], then LFU outperforms LRU. However, if the access pattern has high temporal locality, then LRU can outperform LFU. Therefore, the choice between LRU and LFU depends on the underlying access pattern of the data. This choice, however, does not have to be binary as we can also have an LRFU cache[12] which provides us with a way to have an eviction policy which takes both recency and frequency into account for eviction. This algorithm subsumes both LRU and LFU because its functioning depends on a factor λ which dictates the weightage that is given to recency versus frequency while evicting an item. As a result, we can tune this parameter to have the cache behave like LRU or LFU. Thus, with LRFU cache, we can fine-tune our eviction policy in case we do not have any other additional information about the items.

In the last paragraph we talked about caching policies when we do not have any additional information about the items being stored, however, while processing graphs, we do have some information about the graph topology when we access a node's neighbours. This information can be useful if there is a high likelihood that the neighbours of this node will be accessed. This idea of graph aware caching algorithms was introduced by Aksu et al. in a paper where they introduced a graph aware caching policy named Clock-based graph aware caching (CBGA)[1]. Although the algorithm contain many important contributions, the main idea of this algorithm is to expect that a node's neighbours will be accessed if a node is accessed. This idea helped their algorithm outperform all other competing cache algorithms which did not exploit the structure of the graph. Their idea was further extended by Bok et al.[5] where they added a used-cache in addition to a prefetcher cache for speeding up subgraph lookups. This idea helped them outperform CBGA for the task of subgraph matching where they were able to achieve a hit ratio of between 50% and 65%. These contributions highlight the potential for graph aware caching for reducing the latency of graph access.

Chapter 3

System Architecture

3.1 Component Overview

Figure-3.1 contains the high level overview of the system that we will use for query evaluation. We will now give a brief explanation of the components present in this figure:

1. **LDBC:** The Linked Data Benchmark Council(LDBC) publishes datasets for various types of graph processing workloads. In this thesis, we make use of the Social Network Benchmark (SNB) datasets in order to evaluate the performance of our traversals.
2. **LDBC Converter:** The LDBC dataset are in csv format and contain additional information about nodes and edges. As we will explain in Section-3.3.1, we need the data to be in a particular binary format and we do not have any use for node and edge properties. This component of the system removes the unnecessary information from the LDBC datasets and maps the data into the desired binary format.
3. **AWS S3:** Once the LDBC converter converts the data into a binary format, this data is added to an S3 bucket. A ‘bucket’ in AWS S3 is a container for files.
4. **Graph Access Service:** This service provides an interface for accessing a graph stored in AWS S3. This interface provides the users of this interface with the ability to get a node’s neighbours by providing the source node, edge label, and edge direction. This service contains all the caching mechanisms that we will use to reduce the latency of traversals. We will discuss this service in more detail in Section-3.3.
5. **Graph Algorithm Service:** This service is responsible for using the interface provided by the graph access service to perform traversals and measure their performance. We discuss this component in more detail in Section-3.4.

Looking at the architecture in Figure-3.1, it is natural to wonder why do we have two separate services for accessing the graph and performing traversals on that graph. There are two reasons for this separation: First, it separates the logic for accessing graphs from the logic for performing traversals with minimal overhead; Second, this type of architecture is particularly amenable to creation of a multi-tenant storage layer as used in the Neon database[14]. This is why we have chosen to separate the service responsible for accessing the graph and caching from the service responsible for running traversal algorithms.

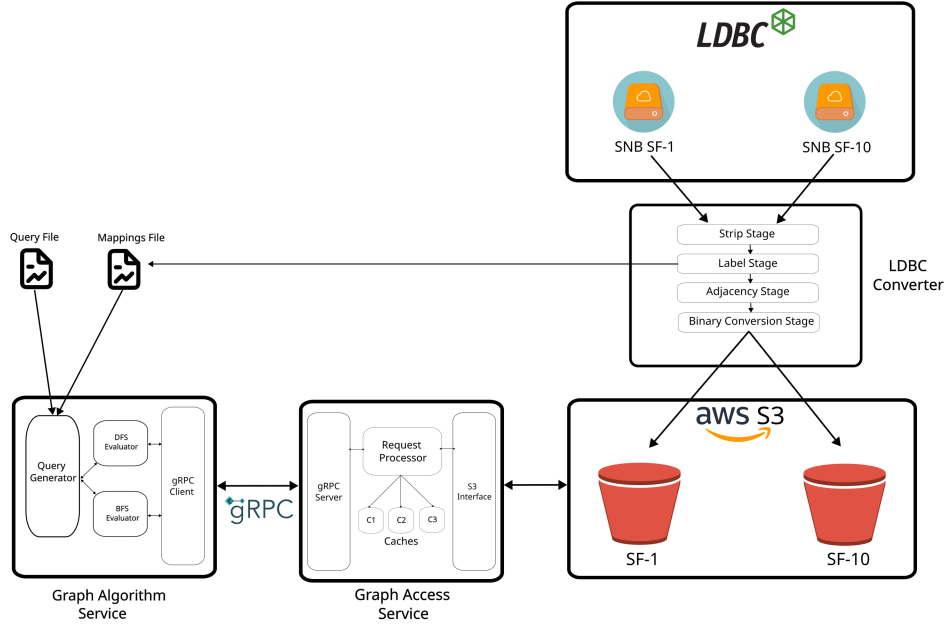


Figure 3.1: High Level System Architecture

3.2 Baseline Implementation

3.3 Graph access service

3.3.1 Modified CSR structure

In the baseline implementation of graph access service described in Section-3.2, we had to download entire files in order to access a node's neighbours. This approach suffers from a couple of problems which makes fetching inefficient and restrictive. The first problem is that we have to fetch and parse the information for an entire file in order to access a single node's neighbours which is wasteful. Furthermore, this approach hinders our ability to have a higher level of granularity for better cache performance. Secondly, this approach makes it harder to utilize the fact that S3 replicates files to various servers. Due to these problems, we decided to modify the underlying binary format and the way we fetch a node's neighbours.

Figure-3.2 shows the binary format that we use for fetching a node's neighbour. This format can logically be divided into three layers. The first layer, which is always of a constant size, contains the first and last node identifier present in this file. The second layer, contains the byte offset for the first incoming edge for a node and the first outgoing edge for a node. Since these offsets are of a constant size and the fact that the first layer of the header tells us how many nodes are present in this file, we can calculate the size of this second layer. The third and final layer contains the edge information, which in our case consists of the edge label and the edge desitnation. This type of structure closely resembles a traditional CSR format except for the fact that we store both incoming and outgoing edges in the same array and the fact that we store byte offsets instead of array indices.

The aforementioned format gives us the ability to fetch a node's neighbours without having to download an entire file. With this modified structure, we can store the first layer for every file in memory and fetch the second layer of a file when it is first accessed. After doing so, we




Figure 3.2: Updated CSR Format

would be able to fetch a node’s incoming or outgoing neighbours without any additional overhead. Furthermore, we can now parallelize requests to a single file which would lead to better throughput since these requests would likely be distributed to different instances within S3. As an additional benefit, we can now perform caching on a more granular level since we do not incur the overhead of fetching and parsing an entire file whenever we need to access a node’s neighbours.

3.3.2 Caching and Prefetching

Figure-3.3 shows the final architecture of the graph access service. The service consists of three basic parts:

1. **Interfaces for external communication:** There are two components which provide an interface with external resources and are colored grey in the figure. The first component, on the left, exposes a gRPC interface which receives requests for accessing a node’s neighbours, starting traversals, and ending traversals. The second component, on the right, is responsible for accessing S3. This component takes a filename along with starting and ending file offsets and returns the requested file content.
2. **Orchestration:** There are two components that help with orchestration of request and are colored blue in the figure. The first component is the ‘Request Processor’ which is responsible for interacting with the caches, moving data between caches, and fetching data from S3 in case it is not present in the caches. The second component titled ‘Vertex offsets’ contains metadata related to the CSR format that we described in Section-3.3.1.
3. **Caches:** There are two different types of caches that are used in this service and are colored green in the figure. The first of these components is an LRFU cache which is used to store nodes that were previously accessed. The second component is a per-query prefetcher which fetches the nodes that are likely to be accessed by a client in the near future. The rest of this section is devoted to the description of these two components.

The **LRFU cache** is a relatively simple




Figure 3.3: Graph Access Service

3.4 Parallelizing graph algorithm service

3.4.1 Parallel BFS implementation

3.4.2 Parallel DFS implementation

Chapter 4

Evaluation

4.1 Comparison with baseline

4.1.1 Optimizing partition sizes

4.2 Comparison with other tools

4.2.1 Neo4j

Monolithic Deployment

Distributed Deployment

4.2.2 Apache Flink

Chapter 5

Discussion

5.1 Alternate system architectures

5.2 Use cases for the system

5.3 Threats to credibility of this work

Chapter 6

Conclusion and Future Work

6.1 Future Work

Bibliography

- [1] Hidayet Aksu, Mustafa Canim, Yuan-Chi Chang, Ibrahim Korpeoglu, and Özgür Ulusoy. Graph aware caching policy for distributed graph stores. In *2015 IEEE International Conference on Cloud Engineering*, pages 6–15. IEEE, 2015. 7
- [2] AWS. Aws aurora serverless. <https://aws.amazon.com/rds/aurora/serverless/>. 5
- [3] AWS. Aws s3. <https://aws.amazon.com/s3/>. 1, 3
- [4] Azure. Azure blob storage. <https://azure.microsoft.com/en-us/products/storage/blobs>. 1
- [5] Kyoungsoo Bok, Seunghun Yoo, Dojin Choi, Jongtae Lim, and Jaesoo Yoo. In-memory caching for enhancing subgraph accessibility. *Applied Sciences*, 10(16):5507, 2020. 7
- [6] Matthias Brantner, Daniela Florescu, David Graf, Donald Kossmann, and Tim Kraska. Building a database on s3. In *Proceedings of the 2008 ACM SIGMOD international conference on Management of data*, pages 251–264, 2008. 5
- [7] Benoit Dageville, Thierry Cruanes, Marcin Zukowski, Vadim Antonov, Artin Avanes, Jon Bock, Jonathan Claybaugh, Daniel Engovatov, Martin Hentschel, Jiansheng Huang, et al. The snowflake elastic data warehouse. In *Proceedings of the 2016 International Conference on Management of Data*, pages 215–226, 2016. 5
- [8] Gil Einziger, Roy Friedman, and Ben Manes. Tinylfu: A highly efficient cache admission policy. *ACM Transactions on Storage (ToS)*, 13(4):1–31, 2017. 6
- [9] Sanjay Ghemawat, Howard Gobioff, and Shun-Tak Leung. The google file system. In *Proceedings of the nineteenth ACM symposium on Operating systems principles*, pages 29–43, 2003. 3
- [10] Tianyang Jiang, Guangyan Zhang, Zican Huang, Xiaosong Ma, Junyu Wei, Zhiyue Li, and Weimin Zheng. {FusionRAID}: Achieving consistent low latency for commodity {SSD} arrays. In *19th USENIX Conference on File and Storage Technologies (FAST 21)*, pages 355–370, 2021. 6
- [11] Douglas J Klein. Centrality measure in graphs. *Journal of mathematical chemistry*, 47(4):1209–1223, 2010. 6
- [12] Donghee Lee, Jongmoo Choi, Jong-Hun Kim, Sam H Noh, Sang Lyul Min, Yookun Cho, and Chong Sang Kim. Lrfu: A spectrum of policies that subsumes the least recently used and least frequently used policies. *IEEE transactions on Computers*, 50(12):1352–1361, 2001. 7
- [13] SUN Microsystems. Networking on the sun workstation, 1986. 3
- [14] Neon. Neon. <https://neon.tech/>. 1, 5, 9
- [15] Google Cloud Platform. Google cloud storage. <https://cloud.google.com/storage>. 1

-
- [16] Siddhartha Sahu, Amine Mhedhbi, Semih Salihoglu, Jimmy Lin, and M Tamer Özsu. The ubiquity of large graphs and surprising challenges of graph processing. *Proceedings of the VLDB Endowment*, 11(4):420–431, 2017. 1
- [17] Konstantin Shvachko, Hairong Kuang, Sanjay Radia, and Robert Chansler. The hadoop distributed file system. In *2010 IEEE 26th symposium on mass storage systems and technologies (MSST)*, pages 1–10. Ieee, 2010. 3
- [18] Snowflake. Snowflake. <https://www.snowflake.com/en/>. 1, 3
- [19] Rebecca Taft, Irfan Sharif, Andrei Matei, Nathan VanBenschoten, Jordan Lewis, Tobias Grieger, Kai Niemi, Andy Woods, Anne Birzin, Raphael Poss, et al. Cockroachdb: The resilient geo-distributed sql database. In *Proceedings of the 2020 ACM SIGMOD international conference on management of data*, pages 1493–1509, 2020. 5
- [20] Alexandre Verbitski, Anurag Gupta, Debanjan Saha, Murali Brahmadesam, Kamal Gupta, Raman Mittal, Sailesh Krishnamurthy, Sandor Maurice, Tengiz Kharatishvili, and Xiaofeng Bao. Amazon aurora: Design considerations for high throughput cloud-native relational databases. In *Proceedings of the 2017 ACM International Conference on Management of Data*, pages 1041–1052, 2017. 5
- [21] George Kingsley Zipf. Relative frequency as a determinant of phonetic change. *Harvard studies in classical philology*, 40:1–95, 1929. 7

Appendix A

System Details

A.1 Query Evaluation

In this file (appendices/main.tex) you can add appendix chapters, just as you did in the thesis.tex file for the ‘normal’ chapters. You can also choose to include everything in this single file, whatever you prefer.